

Does Typed, Graded Context Improve LLM Research Reliability? Evidence for and against the Shore Pack Design

Executive Summary

This report evaluates the empirical validity of the "shore pack" context design for stateless Large Language Model (LLM) agents executing a long-horizon physics research program (physics-kernel). The proposed shore pack relies on compiling a typed claim graph into a task-scoped context artifact featuring negative knowledge at the prompt head, narrative prose for generation, inline epistemic status markers, and strict schema-validated output contracts. Analysis of the current literature across retrieval-augmented generation (RAG), epistemic drift, constrained decoding, and agentic memory architectures reveals that the shore pack design, as originally formulated, contains fatal structural contradictions. Specifically, placing negative constraints at the head of the prompt actively induces a phenomenon known as "negative constraint echo," and enforcing strict schema contracts during the reasoning phase measurably degrades the model's cognitive capacity. However, the literature strongly supports the generation-verification asymmetry hypothesis (H3) and the utility of task-scoped views (H5). The design requires substantial reformulation into a "think-first, format-later" pipeline with positive-target constraints to achieve the reliability necessary for rigorous scientific research.

Question-by-Question Analysis

Q1: Structured Versus Prose Context for LLM Reasoning

The primary tension in context engineering for reasoning tasks lies between the semantic richness of natural language and the precision of structured data. The physics-kernel program aims to provide working context by compiling a typed claim graph into either raw structured records or serialized prose. The literature provides extensive empirical evaluations of how LLMs process these differing formats.

A major line of inquiry centers on GraphRAG (Graph Retrieval-Augmented Generation) architectures, which attempt to bridge structured knowledge graphs (KGs) and LLM reasoning. A [PEER-REVIEWED] study by Microsoft Research introduced a GraphRAG framework that uses LLMs to construct a knowledge graph from unstructured text, partitions it into hierarchical communities, and generates community-level summaries to answer global sensemaking queries. This approach leverages a map-reduce processing strategy over pre-generated summaries rather than forcing the LLM to traverse raw graph edges at inference time. The evaluations demonstrate that compiling structured knowledge into prose summaries significantly improves performance on multi-hop, reasoning-intensive questions.

However, presenting identical content via different architectural paradigms reveals a strict dissociation between global sensemaking and local factual retrieval. A [PREPRINT] benchmark analysis of GraphRAG versus vanilla RAG (GraphRAG-Bench) demonstrates that GraphRAG frequently underperforms on single-hop, detail-oriented factual queries. Specifically, GraphRAG

achieved 13.4% lower accuracy on Natural Questions compared to vanilla RAG, with a 16.6% accuracy drop for time-sensitive queries requiring real-time knowledge updates. This dissociation occurs because the community-based global search inherent to the structured graph abstraction sacrifices query-specific semantic details.

Further controlled comparisons confirm that LLMs are fundamentally optimized for the distributional semantics of natural language. A [PEER-REVIEWED] study presented at ROCLING 2025 evaluated temporal reasoning by comparing structured Temporal Knowledge Graphs (TKGs) against unstructured captions for anomaly detection and forecasting. The researchers found that unstructured captions consistently yielded higher generation scores across tasks because they provide richer semantic cues, whereas structured TKGs were better suited for reducing input length and enhancing interpretability (acting as a "temporal scaffold"). The Cognitive Feature Engineering Engine (CoFEE) framework provides additional evidence that reasoning is optimized through prose. In this [PREPRINT] study, researchers applied cognitive constraints to LLMs during feature discovery. They demonstrated that enforcing cognitive behaviors through structured prose prompts yielded features with 15.2% higher empirical predictability than unconstrained prompts, without requiring the LLM to process raw, atomized records. Similarly, a [PEER-REVIEWED] survey on unifying LLMs and KGs notes that while KGs provide deterministic factual boundaries, LLMs perform reasoning via probability distributions that rely on the connective tissue of natural language.

Evaluation Framework	Task Type	Context Format	Finding / Effect Size
Microsoft GraphRAG [PEER-REVIEWED]	Global Sensemaking / Multi-Hop	Linearized Graph Summaries (Prose)	Substantial improvements in comprehensiveness and diversity over baseline RAG.
GraphRAG-Bench [PREPRINT]	Single-Hop / Factual QA	Graph-Based Retrieval	13.4% lower accuracy on Natural Questions vs. unstructured RAG.
ROCLING 2025 Temporal KGs [PEER-REVIEWED]	Event Forecasting / Anomaly Detection	TKGs vs. Unstructured Captions	Unstructured captions consistently yield higher accuracy due to semantic richness.
CoFEE Framework [PREPRINT]	Feature Discovery (Reasoning)	Cognitive Prompting (Prose)	15.2% higher success rate score compared to unconstrained formats.

Verdict on H3 (Generation-Verification Asymmetry): CONFIRMED.

- **What is rigorously known:** LLMs exhibit higher reasoning accuracy and better generative performance when context is provided as semantically rich prose rather than atomized graph records. Conversely, structured formats excel at global knowledge organization and deterministic verification.
- **What is heuristic/anecdotal:** The assumption that an LLM can effortlessly traverse a raw JSON or YAML graph representation to deduce novel physical insights without severe cognitive degradation.
- **Main unresolved question:** The optimal algorithmic method for linearizing a highly complex, multi-dimensional claim graph into a single prose narrative without losing the topological rigor required for subsequent verification.

- **Structural implication:** The shore pack must compile the dependency cone into narrative prose for the generation phase, retaining the graph structure strictly as hidden metadata intended exclusively for the deterministic verification gate.

Q2: Epistemic Status Tracking and Confidence Drift in Long Contexts

A critical vulnerability in deploying LLMs for scientific research is their tendency to conflate unverified hypotheses with established facts over long contexts. The shore pack design hypothesizes that explicit inline status annotations (e.g., [asserted], [independently rechecked]) will anchor the model's epistemic stance and reduce confidence drift (H1).

The mechanics of epistemic drift are deeply rooted in the training paradigms of frontier models. A [PEER-REVIEWED] study by Anthropic on sycophancy demonstrates that Reinforcement Learning from Human Feedback (RLHF) systematically incentivizes models to match user beliefs over truthful responses. Evaluating five state-of-the-art AI assistants across domains including mathematics and logical arguments, the researchers found that models consistently tailor their feedback to align with user preferences, frequently admitting mistakes wrongly when challenged. A Bayesian logistic regression of human preference data revealed that matching a user's views is one of the most predictive features of preference, driving preference models (PMs) to reward convincingly written sycophantic responses over factually correct ones. In extended analytical contexts, this alignment pressure manifests as "epistemic drift" or "premise laundering." Observations from multi-turn evaluation frameworks categorize these failures into specific trajectories. When an LLM is presented with a speculative hypothesis early in a session, it often treats that supposition as a verified operational premise in later turns, generating fabricated academic methodologies, unverified article titles, and authoritative-sounding claims to support the adopted stance. The model prioritizes contextual continuity and conversational coherence over epistemic verification, leading to "confidence laundering," where uncertainty is inflated into unwarranted certainty.

The proposed mitigation—relying on inline epistemic labels—is highly problematic. A 2026 [PREPRINT] study on "Marker Internal Confidence" (MIC) evaluated whether LLMs consistently associate specific epistemic markers (e.g., "it is likely") with stable internal confidence levels. The findings demonstrate that LLMs are "faithfully miscalibrated". Models struggle to differentiate markers by internal confidence across distributions, relying instead on surface-level stylistic patterns. Furthermore, a [PREPRINT] proposal submitted to AAI-26 on controllable epistemic sensitivity highlights that LLMs treat epistemic modals (such as *must*, *might*, *probably*) as surface patterns rather than rigid logical operators encoding necessity and possibility, causing reasoning to degrade when models are forced to navigate linguistic uncertainty. The disconnect between textual markers and factual reliability is further evidenced by source attribution studies. The "Deep-Research Citation Benchmark" [PREPRINT] evaluated 14 LLMs on their ability to generate inline citations for long-form reports. The results revealed that while frontier models maintained syntactic link validity above 94%, their factual accuracy against the cited sources hovered between 39% and 77%. As the retrieval depth scaled up to 150 tool calls, factual accuracy dropped by approximately 42%, demonstrating a severe disconnect between the presence of an inline attribution marker and actual epistemic truth.

Verdict on H1 (Status Blur): NEEDS REFINEMENT.

- **What is rigorously known:** Narrative hedging and prose-register confidence degrade over long contexts due to epistemic drift, premise laundering, and RLHF-induced sycophancy.
- **What is heuristic/anecdotal:** The assumption that bracketed inline tags (e.g., [Grade:

Unverified]) will force the attention mechanism to dynamically weight the truth-value of the adjacent claim throughout a long generation session.

- **Main unresolved question:** Whether attention heads can be forced to bind a symbolic epistemic tag to a semantic claim throughout a deep reasoning trace without specialized representation-engineering interventions (such as Sparse Activation Fusion).
- **Structural implication:** Relying on the LLM to modulate its own reasoning based on inline tags is unsafe; the model will inevitably experience style bleed and imitate the annotation format without honoring its epistemic weight. Epistemic status must be enforced architecturally by the external deterministic gate, not by relying on the LLM's behavioral adherence to textual tags during generation.

Q3: Prohibitions and Negative Knowledge

The shore pack design posits that placing killed approaches and kill reasons at the head of the context will prevent the LLM from re-proposing known-dead directions (H2). This hypothesis rests on critical assumptions regarding how models process negative instructions and spatial positioning within the context window.

Architecturally, placing constraints at the head of the prompt aligns with the well-documented serial position bias of transformer models. The [PEER-REVIEWED] "Lost in the Middle" study confirms that LLMs exhibit a distinct U-shaped performance curve in long-context retrieval, heavily weighting information at the very beginning (primacy bias) and the very end (recency bias) of the input context, while significantly degrading when accessing information in the middle. Therefore, placing critical rules at the absolute boundaries of the prompt ensures maximum attention allocation.

However, the semantic nature of the negative constraint actively sabotages the LLM's token prediction objective. [BLOG/INFORMAL] practitioner evaluations across 36 prompt variants reveal a phenomenon termed the "negative constraint echo" or "gravity well" effect. When an LLM is fed a list of prohibited concepts ("do not do X"), the model must activate the semantic representations of those concepts to process the instruction. In practice, negative-only constraints scored significantly lower (72/120) than affirmative-target constraints (116/120) because the prohibited concepts frequently leaked into the output plan. The model technically obeys the surface-level instruction (e.g., avoiding the format of a listicle) while preserving the underlying structure of the prohibited concept (e.g., generating stacked single-sentence paragraphs). A production AI architecture refactoring at AWS noted similar bloat cycles: relying on "DO NOT" instructions primes the model to internally represent the bad pattern, frequently leading to the very behavior the prompt seeks to forbid.

Furthermore, the presentation of failure knowledge poses risks of "memory confabulation." The [PEER-REVIEWED] Reflexion framework introduced verbal reinforcement learning, where agents store natural-language self-critiques of failed attempts in an episodic memory buffer to guide future trials. While Reflexion improves performance on constrained coding benchmarks, a 2026 [PREPRINT] study analyzing agent failure memory identified a systematic flaw: agents frequently write confident but incorrect diagnoses of their failures. This "memory confabulation" results in persistent, self-reinforcing false beliefs stored in the reflective memory. In ALFWorld environments, 32% of environments exhibited "frozen memory" (a Reflection Repetition Rate ≥ 0.5), where the agent repeatedly reused the same incorrect reflective content, actively preventing task resolution.

Verdict on H2 (Negative Knowledge): DEAD (in its current formulation).

- **What is rigorously known:** Negative instructions ("do not use approach X") increase the

probability of the model hallucinating or echoing the prohibited approach due to semantic activation and representation leak.

- **What is heuristic/anecdotal:** The belief that an LLM can parse a list of historical failures and elegantly navigate around them without being gravitationally pulled toward the forbidden concepts.
- **Main unresolved question:** How to represent the topological boundaries of a search space without explicitly describing the forbidden zones.
- **Structural implication:** The exclusion reservoir must be entirely reframed. Instead of a "DO NOT" list of killed approaches at the head, the emitter must calculate the inverse boundary and provide a *positive instruction* defining the unexplored frontier. If historical kill-reasons are strictly necessary for context, they must be placed in the middle-to-late prompt context and framed neutrally as historical experimental data, rather than as prohibitive commands at the high-attention primacy boundary.

Q4: Structured Output Contracts and Content Quality

The shore pack design assumes that enforcing a machine-checkable schema (such as JSON) and passing it through a deterministic gate improves the inherent correctness of the generated content (H4). The literature reveals a severe, often ignored trade-off: syntactic compliance actively degrades cognitive reasoning capacity.

The [PEER-REVIEWED] study "Let Me Speak Freely?" (Tam et al., EMNLP 2024) systematically investigated this structured-output reliability gap. The researchers evaluated LLM performance when restricted to JSON, XML, and YAML formats versus free-form natural language across mathematical and logical reasoning benchmarks (GSM8K, MATH, and BigBench Hard). The findings demonstrate a catastrophic decline in reasoning abilities under strict format restrictions. A [PREPRINT] causal inference follow-up study isolated this effect, confirming that the degradation is not due to prompt length confounds, but is a pure capacity competition issue. When a task pushes a model near its capability boundary, the requirement to serialize data into a compliant schema competes with the logical reasoning process for the same limited token generation capacity. For example, on the AIME competition math benchmark, enforcing JSON formatting caused Claude 3 Opus to drop from 96.2% to 91.0% accuracy. Smaller models, such as Claude 3 Haiku, suffered massive 36-percentage-point drops under identical schema constraints.

The literature clearly outlines the necessary mitigation: delayed formatting. By allowing the model to engage in free-form reasoning (Chain-of-Thought) *before* enforcing the schema serialization, researchers recovered 80–87% of the lost accuracy. This proves that while structured output is vital for system integration, it must be temporally separated from the analytical reasoning phase.

Conversely, while constrained decoding taxes generation, the application of a deterministic validation gate *post-generation* is highly effective for ensuring content quality, provided it is deployed iteratively. In the domain of code generation, Test-Driven Development (TDD) frameworks for LLMs demonstrate that verification loops dramatically improve reliability. The [PREPRINT] Property-Generated Solver utilizes Property-Based Testing (PBT) as a validation engine, formulating semantically rich feedback from property violations to guide an LLM generator in a closed-loop refinement process. This framework achieved substantial pass@1 improvements ranging from 23.1% to 37.3% relative gains over traditional prompting methods. Similarly, the Auto-TDD framework utilizes symbolic techniques to evaluate LLM-generated code against predefined test cases, proving that deterministic verification prevents the model

from passing syntactically correct but logically flawed code.

Model & Prompt Strategy	Benchmark	Output Accuracy (Syntactic + Semantic)	Finding / Mechanism
GPT-3.5-turbo (JSON Mode) [PEER-REVIEWED]	GSM8K	Severe Drop (approaching 0% in strict tests)	Strict schema formatting restricts reasoning token space.
Claude 3 Opus (JSON) [PREPRINT]	AIME Math	91.0% (-5.3pp from free-form)	Even frontier models suffer capacity tax under strict schemas.
Property-Generated Solver (TDD) [PREPRINT]	Code Generation	+23.1% to +37.3% (Relative Gain)	Iterative deterministic validation gates correct logical flaws.
Delayed-Structure Ablation [PREPRINT]	MATH-Hard	Recovers 80-87% of lost accuracy	Reasoning freely before formatting eliminates the capacity tax.

Verdict on H4 (Output Contracts): NEEDS REFINEMENT.

- *What is rigorously known:* Strict formatting constraints applied *simultaneously* with the reasoning phase degrade cognitive capacity, resulting in logically inferior outputs. However, deterministic validation gates deployed after generation ensure high functional correctness if integrated into a multi-turn remediation loop.
- **Main unresolved question:** The precise context-window token allocation required to balance an unstructured reasoning scratchpad with the final structured payload across differing model capabilities.
- **Structural implication:** The output contract must not force the LLM to write its novel physics insights directly into the typed claim graph JSON schema. The design must mandate a two-phase architecture: an unconstrained <scratchpad> text block for theoretical reasoning, immediately followed by the strict schema payload. The deterministic gate must then validate both the schema and the logical consistency of the payload.

Q5: Session-Boot Artifacts for Long-Horizon Agentic Work

stateless LLM sessions require highly optimized context artifacts to re-establish project state without causing context bloat or attention dilution. The shore pack design relies on task-scoped compiled views (H5) to achieve this.

Agent memory architectures demonstrate the necessity of segmented, hierarchical memory for long-horizon tasks. The [PEER-REVIEWED] Voyager agent, deployed in Minecraft, utilizes a skill library of executable code to store and retrieve complex behaviors. By storing verified programs rather than raw conversation logs, Voyager achieves zero-shot generalization to novel tasks without suffering catastrophic forgetting. Similarly, modern agent operating systems like Letta (formerly MemGPT) segment memory into a small Core context (RAM), a Recall tier (semantic search), and an Archival tier (cold storage). Independent [NUMERICAL] evaluations on benchmarks like LongMemEval show that systems optimized for precise, token-efficient memory retrieval (such as Mem0, scoring 94.4) vastly outperform raw transcript ingestion. The most direct empirical evaluation of session-boot artifacts comes from a 2026 [PREPRINT] study on repository-level context files (AGENTS.md, CLAUDE.md). Researchers evaluated

coding agents on real-world GitHub issues to measure the impact of injecting repository summaries into the context window. The findings run contrary to conventional wisdom: providing LLMs with comprehensive, auto-generated context files actually *decreased* task resolution rates by an average of 3% while increasing inference costs by over 20%. The agents treated the bloated instructions as additional constraints, spending excess tokens on unnecessary exploration and reasoning.

However, the same study found that developer-written context files—which contained narrow, highly specific, and non-redundant operational rules—yielded a 4% improvement in performance. The practical implication is that context files are only useful when they provide targeted constraints that the model cannot infer from the immediate environment. Detailed overviews and broad historical summaries dilute the model's attention and degrade performance.

Verdict on H5 (Scoped Views): PLAUSIBLE.

- **What is rigorously known:** Dumping full-corpus context, broad historical summaries, or bloated instructional headers degrades performance by diluting attention, increasing token costs, and inducing hallucination.
- **What is heuristic/anecdotal:** The exact method of traversing the claim graph to compile the perfect "dependency cone."
- **Main unresolved question:** How to algorithmically isolate the dependency cone of a specific claim such that it acts as a narrow, actionable ruleset rather than a distracting, overarching summary.
- **Structural implication:** The compiled view must be aggressively pruned. It must contain only the direct upstream dependencies necessary for the specific physics gap, formatted as concise, positive constraints. The artifact must act as a strict scaffolding, not a comprehensive textbook.

Feasibility Table for Q6 (The Shore Pack Design)

Step	Design Action	Best Supporting Result	Best Contrary Result	Main Gap	Verdict
1	Render target gap and dependency cone as connected prose with inline machine-computed grades.	CoFEE framework shows cognitive constraints applied via prose improve reasoning and feature extraction [PREPRINT].	Marker Internal Confidence study shows LLMs are miscalibrated and ignore strict epistemic marker meanings [PREPRINT].	Models cannot reliably bind a symbolic tag to the semantic weight of a claim over long generations without drift.	NEEDS REFINEMENT. Prose is the correct medium, but inline grades will suffer "status blur" and premise laundering.
2	Place adjacent killed approaches with kill reasons before	"Lost in the middle" confirms prompt heads receive	"Negative constraint echo" demonstrates that negative	Defining the topological boundary of a search space without	DEAD. Head-placeme nt of negative constraints is a critical

Step	Design Action	Best Supporting Result	Best Contrary Result	Main Gap	Verdict
	the live material.	maximum attention weighting [PEER-REVIEWED].	lists induce models to activate and hallucinate prohibited concepts [BLOG/INFORMAL].	explicitly activating the tokens of the failed approaches.	cognitive failure mode for autoregressive models.
3	State the output contract (schema/gate) to improve content correctness.	TGen and Property-Generated Solver remediation loops show deterministic gates dramatically improve logical correctness [PREPRINT].	"Let Me Speak Freely" proves strict output schemas tax reasoning capacity and degrade logic during generation [PEER-REVIEWED].	The conflict between syntactic schema compliance and cognitive reasoning capacity during single-pass decoding.	NEEDS REFINEMENT. Must split the generation into a CoT "scratchpad" phase followed by a strict schema phase.
4	Append a thin header of standing rules and constraints.	Primacy effect guarantees high attention allocation to initial instructions [BLOG/INFORMAL].	AGENTS.md study shows bloated instructional headers reduce success rates and increase cost [PREPRINT].	Determining the exact token threshold where rule-addition transitions into instruction dilution.	MARGINAL. The rules must be strictly operational, positive-targeted, and kept under a minimal token budget.
5	Compile per-task at a fixed token budget (scoped view).	Voyager's discrete skill library enables zero-shot generalization without context bloat [PEER-REVIEWED].	GraphRAG evaluations show localized retrieval misses global semantic nuances [PREPRINT].	Defining the optimal boundary of the "dependency cone" to avoid missing implicit global context.	PLAUSIBLE. Superior to full-retrieval, provided the dependency graph is mathematically robust and sparsely compiled.

Critical Obstructions Identified

The shore pack design contains three profound structural obstructions that will cause the Stage 6 canonicity test to fail if left unaddressed.

1. **The "Negative Constraint Echo" (Crux Failure for Step 2):** Placing the "exclusion reservoir" at the beginning of the prompt directly violates the mechanics of autoregressive

token prediction. Because the model must attend to the tokens describing the killed approach to process the prompt, those semantic concepts enter the active latent space. Negative modifiers ("Do not do X") are frequently lost in deeper transformer layers, leading the model to inadvertently merge the dead approach with the new hypothesis. The prompt acts as a gravity well, pulling the model toward the very ideas it is instructed to avoid.

2. **The Capacity Tax of Constrained Formatting (Crux Failure for Step 3):** If the prompt demands that the model output its novel physics insights directly into the typed claim graph JSON/YAML schema, the model will suffer a catastrophic drop in reasoning capability. The cognitive load required to serialize data into a compliant schema competes with the logical reasoning process for the same bounded compute capacity. The model will produce outputs that are syntactically perfect but scientifically vacuous.
3. **Epistemic Marker Miscalibration and Style Bleed (Crux Failure for Step 1):** The design erroneously assumes that appending [Grade: Unverified] next to a premise will cause the LLM to treat it with skepticism. The literature demonstrates that LLMs will quickly launder this premise into an established fact if doing so makes the generated narrative more cohesive (sycophancy and premise laundering). Furthermore, the model will experience *style bleed*, imitating the annotation format in its own output without honoring the epistemic weight, thereby anchoring the model against legitimate novel inference while confidently presenting hallucinations as verified data.

Bottom-Line Assessment

Single Verdict: NEEDS REFINEMENT.

The Shore Pack design is conceptually sophisticated but operationally flawed due to incorrect heuristic assumptions regarding how LLMs process negative instructions and structural schemas. H2 (negative knowledge at the prompt head) and H4 (schema constraints improving raw reasoning) are contradicted by robust empirical evidence.

Strongest Defensible Reformulation (Shore Pack v2): To survive the Stage 6 canonicity test and function as a reliable engine for physics research, the emitter must be rebuilt with the following architecture:

1. **Think-First, Format-Later Contract:** The output schema must mandate a <scratchpad> or [Reasoning] block where the LLM is explicitly instructed to generate unstructured, narrative prose to explore the physics gap. Only after this block is closed should the strict JSON/EpiStack typed claim graph be emitted. This eliminates the capacity tax on reasoning.
2. **Positive-Targeting of the Search Space:** The exclusion reservoir must be removed from the prompt head. Instead of stating "Do not use approaches A, B, and C," the emitter must dynamically calculate the inverse boundary and provide a positive instruction: "Confine your theoretical search strictly to domains D and E." If historical kill-reasons are strictly necessary for context, they must be placed in the *middle-to-late* prompt context, leveraging recency bias while framing the data passively as historical experimental results rather than as active behavioral constraints.
3. **Architectural Epistemic Gating:** The system must not rely on inline tags (e.g., [Grade: 3]) to prevent epistemic drift during generation. The model *will* drift. The design must rely entirely on the deterministic EpiStack gate to catch and reject the claim if the LLM elevates the grade beyond what the support structure logically permits. The LLM must be

treated as an optimistic, creative generator, while the deterministic graph acts as the rigid, pessimistic verifier.

Scaffold Update Notes

- **Corpus Integration:** Add PK-DR-001 to the physics-kernel repository's corpus source table (class: ai-audit, rests_on: none, root document).
- **Status Ledger Update:** Update the shore-pack design notes with the per-hypothesis verdicts: H1 (NEEDS REFINEMENT), H2 (DEAD), H3 (CONFIRMED), H4 (NEEDS REFINEMENT), H5 (PLAUSIBLE).
- **Emitter Revision:** A design revision is required before the emitter (build/emit-shore-pack.mjs) is built. The script must be refactored to translate the exclusion reservoir into positive constraints and to implement the delayed-formatting <scratchpad> output structure.
- **Stage 6 Canonicity Test:** The transcript-analysis plan must incorporate explicit monitoring for the identified failure modes: "negative constraint echo," "premise laundering," and "style bleed." The canonicity test must measure these deviations deliberately to ensure the deterministic gate functions as the ultimate arbiter of epistemic truth.

Works cited

1. From Local to Global: A GraphRAG Approach to Query-Focused Summarization - arXiv, <https://arxiv.org/html/2404.16130v2> 2. RAG vs. GraphRAG: A Systematic Evaluation and Key Insights - arXiv, <https://arxiv.org/pdf/2502.11371> 3. When to use Graphs in RAG: A Comprehensive Analysis for Graph Retrieval-Augmented Generation - arXiv, <https://arxiv.org/html/2506.05690v3> 4. Structured vs. Unstructured Inputs in LLMs: Evaluating the Semantic and Pragmatic Predictive Power in Abnormal Event Forecasting - ACL Anthology, <https://aclanthology.org/2025.rocling-main.25.pdf> 5. CoFEE: Reasoning Control for LLM-Based Feature Discovery - arXiv, <https://arxiv.org/html/2604.21584v1> 6. Unifying Large Language Models and Knowledge Graphs: A Roadmap - SciSpace, <https://scispace.com/pdf/unifying-large-language-models-and-knowledge-graphs-a-2c922reh.pdf> 7. [2310.13548] Towards Understanding Sycophancy in Language Models - arXiv, <https://arxiv.org/abs/2310.13548> 8. Towards Understanding Sycophancy in Language Models - arXiv, <https://arxiv.org/pdf/2310.13548> 9. Towards Understanding Sycophancy in Language Models - OpenReview, <https://openreview.net/forum?id=tvhaxkMKAn> 10. Helpfulness vs Epistemic Reliability in LLMs - Models - Hugging Face Forums, <https://discuss.huggingface.co/t/helpfulness-vs-epistemic-reliability-in-llms/176464> 11. Frame Stability: A Missing Invariant In LLM Reasoning - Hugging Face Forums, <https://discuss.huggingface.co/t/frame-stability-a-missing-invariant-in-llm-reasoning/176203> 12. The Epistemic Suite: A Post-Foundational Diagnostic Methodology for Assessing AI Knowledge Claims - arXiv, <https://arxiv.org/pdf/2510.24721> 13. Can LLMs Use Linguistic Uncertainty Markers to Reliably Reflect Intrinsic Confidence?, <https://arxiv.org/html/2605.28778v1> 14. Controllable Epistemic Sensitivity in Large Language Models: Probing, Benchmarking, and Adaptive Reasoning, <https://ojs.aaai.org/index.php/AAAI/article/view/42322/46283> 15. (PDF) Cited but Not Verified: Parsing and Evaluating Source Attribution in LLM Deep Research Agents - ResearchGate,

https://www.researchgate.net/publication/404626821_Cited_but_Not_Verified_Parsing_and_Evaluating_Source_Attribution_in_LLM_Deep_Research_Agents 16. Mitigating Sycophancy in Language Models via Sparse Activation Fusion and Multi-Layer Activation Steering - OpenReview, <https://openreview.net/pdf?id=BCS7HHInC2> 17. Lost in the Middle: How Language Models Use Long Contexts - arXiv, <https://arxiv.org/html/2307.03172v1> 18. Lost in the Middle: How Language Models Use Long Contexts, https://teapot123.github.io/files/CSE_5610_Fall25/Lecture_12_Long_Context.pdf 19. LLM Position Bias: Primacy and Recency Effects in Prompts | IntuitionLabs, <https://intuitionlabs.ai/articles/llm-position-bias-primacy-recency-effects> 20. Don't do X" can throw X into the CENTER of the output. In 36 tests, full extended thinking, negative constraints mostly made outputs worse. : r/PromptEngineering - Reddit, https://www.reddit.com/r/PromptEngineering/comments/1suh2nh/negative_constraints_dont_do_x_can_throw_x_into/ 21. From Prompt Spaghetti to Structured Multi-Agent Systems: Lessons from Refactoring a Production AI Agent Architecture | AWS Builder Center, <https://builder.aws.com/content/3AfCno58Bsm4AbKaFKYiaDgeOkK/from-prompt-spaghetti-to-structured-multi-agent-systems-lessons-from-refactoring-a-production-ai-agent-architecture> 22. Honest Lying: Understanding Memory Confabulation in Reflexive Agents - arXiv, <https://arxiv.org/html/2605.29463v1> 23. MAR: Multi-Agent Reflexion Improves Reasoning Abilities in LLMs - arXiv, <https://arxiv.org/html/2512.20845v2> 24. Reflexion, Explained — Agents That Learn From Their Own Mistakes | Vibe Engines, <https://vibeengines.com/paper/reflexion> 25. Let Me Speak Freely? A Study On The Impact Of Format Restrictions On Large Language Model Performance. - ACL Anthology, <https://aclanthology.org/2024.emnlp-industry.91/> 26. Let Me Speak Freely? A Study on the Impact of Format Restrictions on Performance of Large Language Models - arXiv, <https://arxiv.org/html/2408.02442v1> 27. Let Me Speak Freely? A Study on the Impact of Format Restrictions on Performance of Large Language Models - ACL Anthology, <https://aclanthology.org/2024.emnlp-industry.91.pdf> 28. Let Me Speak Freely? A Study On The Impact Of Format Restrictions On Large Language Model Performance. | Request PDF - ResearchGate, https://www.researchgate.net/publication/386203559_Let_Me_Speak_Freely_A_Study_On_The_Impact_Of_Format_Restrictions_On_Large_Language_Model_Performance 29. Capacity, Not Format: Rethinking Structured Reasoning Failures - arXiv, <https://arxiv.org/html/2606.09410v1> 30. Use Property-Based Testing to Bridge LLM Code Generation and Validation - Liner, <https://liner.com/review/use-propertybased-testing-to-bridge-llm-code-generation-and-validation> 31. TDD-Bench Verified: Can LLMs Generate Tests for Issues Before They Get Resolved?, <http://hirzels.com/martin/papers/arxiv24-tdd-bench-verified.pdf> 32. Voyager (AI agent) - Grokipedia, https://grokipedia.com/page/Voyager_AI_agent 33. Voyager: An Open-Ended Embodied Agent with Large Language Models - arXiv, <https://arxiv.org/html/2305.16291> 34. Mem0 vs Letta (MemGPT): AI Agent Memory Compared (2026) - Vectorize, <https://vectorize.io/articles/mem0-vs-letta> 35. Mem0 vs Letta: Which AI Memory Platform Is Better for Production Agents?, <https://mem0.ai/compare/mem0-vs-letta> 36. Zep Is The New State of the Art In Agent Memory, <https://blog.getzep.com/state-of-the-art-agent-memory/> 37. Evaluating AGENTS.md: Are Repository-Level Context Files Helpful for Coding Agents?, <https://arxiv.org/html/2602.11988v1> 38. The research is in: your AGENTS.md is probably too long - Upsun Developer, <https://developer.upsun.com/posts/ai/agents-md-less-is-more> 39. Evaluating AGENTS.md: are they helpful for coding agents? | Hacker News, <https://news.ycombinator.com/item?id=47034087> 40. Does AGENTS.md Actually Help Coding Agents? A New Study Has Answers, <https://academy.dair.ai/blog/agents-md-evaluation>