Given the Kernel, What Is a Physics Synthesis? Precedents from How Physics Actually Integrates Its Knowledge

1. Executive Summary

The physics-synthesis affordance layer seeks to harvest the institutional, mathematical, and sociological machinery that the physics community has spent a century refining. By translating this machinery into the strict invariants of the physics-kernel (PK) typed claim graph, the architecture can mechanize the handling of correlated measurements, contradictory results, graded evidence, and dead hypotheses. This report dissects the mechanisms of data-side and theory-side synthesis across canonical physics institutions—including the Particle Data Group (PDG), CODATA, the Flavour Lattice Averaging Group (FLAG), nuclear data evaluation (ENDF), and the Swampland program.

The analysis confirms that the kernel's structural invariants (I1 through I9) perfectly map onto the manual oversight functions currently performed by physics review committees. The three strongest physics-native affordances identified for immediate transfer into the stateless LLM session architecture are:

1. **The No-Go Registry (Mapped to Invariant I3):** The systematic tracking of first-class kills alongside explicit, typed evasion conditions. Physics treats no-go theorems not as the end of inquiry, but as precise topological boundaries. This directly informs the kernel's exclusion reservoir.
2. **FLAG-Style Graded Entry Criteria (Mapped to Invariant I1):** The rigorous, color-coded grading rubric utilized by lattice QCD to determine which theoretical calculations are permitted to enter a synthesis. This provides a mechanical precedent for the kernel's typed, graded claims, proving that theory-side synthesis requires explicit metadata tagging of derivations.
3. **The Tension Dashboard (Mapped to Invariants I5 and I6):** Modeled on the management of the muon anomalous magnetic moment and the proton radius puzzle, this affordance holds contradictions as load-bearing structural features rather than fatal errors. It leverages provenance tracking to compute cruxes and isolate contamination cones originating from shared unstated assumptions.

The Central Disanalogy: Every precedent institution evaluated in this report synthesizes across an adversarial, distributed community of independent research groups [PEER-REVIEWED]. The physics-kernel, however, synthesizes one research program's internal state under a single human director utilizing stateless LLM sessions. This difference is irrelevant for topological tasks—such as tracking the dependency graph of a no-go theorem or enforcing a grading rubric. It is, however, fatal for statistical procedures that assume independent systemic errors. For instance, applying a PDG-style scale factor to expand uncertainties assumes random, uncorrelated biases between independent laboratories [PEER-REVIEWED]. In a stateless LLM architecture, systemic errors (e.g., algorithmic biases or shared pre-training data

hallucinations) are highly correlated. Consequently, statistical averaging mechanisms must be deprecated in favor of logical crux computations.

2. Data-Side Machinery: Dissecting Canonical Institutions

Physics has developed rigorous mathematical and procedural frameworks for combining inconsistent data sets. The objective of these institutions is rarely to merely find a mathematical mean; rather, it is to quantify the reliability of the underlying tension and establish strict rules for inclusion, contradiction, and retirement.

2.1 The Particle Data Group (PDG): Averages and the Scale Factor

The Particle Data Group (PDG) produces the *Review of Particle Physics*, maintaining a vast database of particle properties, gauge bosons, leptons, quarks, mesons, and baryons [PEER-REVIEWED]. The core mechanism for data-side synthesis is the inverse-variance weighted average, expressed as $\hat{\mu} = \sum (x_i/\sigma_i^2) / \sum (1/\sigma_i^2)$, which employs the inverse of the square of the associated uncertainties as weights [PEER-REVIEWED].

However, particle physics data frequently features a spread significantly larger than individual reported uncertainties, indicating uncontrolled systematic effects or biases between different measurement techniques [PEER-REVIEWED]. To synthesize over these held contradictions, the PDG employs a rigid procedural rule based on the reduced chi-squared statistic, defined as $\chi^2/(N-1)$, where N is the number of data points [PEER-REVIEWED]. The procedure dictates:

1. If the reduced χ^2 is less than or equal to 1, the standard weighted average and its uncertainty are accepted without modification.
2. If the reduced χ^2 is greater than 1, indicating internal inconsistency among the aggregated measurements, the PDG artificially enlarges the uncertainty of the average. This is achieved by applying a common scale factor, $S = \sqrt{\chi^2 / N_{eff}}$, to all individual errors before recomputing the final uncertainty [PEER-REVIEWED].
3. If some individual errors are vastly smaller than others, the S factor is computed solely from the most precise experiments to prevent highly uncertain outliers from inflating the scale factor unjustifiably [PEER-REVIEWED].

To visually and analytically represent these tensions, the PDG generates "ideograms," which superimpose individual Gaussian distributions of the data points. This highlights outlier clusters and structural tensions, visually separating data that pull the average in contradictory directions [PEER-REVIEWED]. Regarding retirement rules, deprecated results are handled across editions by moving them to obsolete archives; historically, the PDG maintains a roughly 60:40 ratio of data accepted for averaging versus lesser-quality data that is encoded but excluded [PEER-REVIEWED]. Measurements are excluded if superseded by a vastly more precise methodology or if a fundamental paradigm shift renders the measurement parameters invalid.

Mapping and Disanalogy: The PDG scale factor is a procedural mechanism for managing Invariant I6 (held contradictions with computed cruxes). It allows the synthesis to output a downgraded but usable value rather than failing entirely. The structural property doing the work is the penalization of confidence strictly proportional to the variance of the contradiction. The fatal disanalogy is that the Birge ratio and the S factor assume independent inter-laboratory errors [PEER-REVIEWED]. In the physics-kernel, contradictions arise from divergent logical branches generated by an LLM, not independent sensor noise. Therefore, scaling an error bar mathematically is inappropriate; the kernel-native analogue must utilize I6 to generate a qualitative crux rather than a quantitative smoothing factor.

2.2 CODATA: Least-Squares Adjustments and Discrepancy Expansion

The Committee on Data of the International Science Council (CODATA) periodically adjusts the recommended values of fundamental physical constants [PEER-REVIEWED]. The machinery relies on a multivariable least-squares adjustment (LSA) that minimizes the weighted sum of squared residuals, operating over complex covariance matrices [PEER-REVIEWED]. The LSA solves an overdetermined system of equations $Ax = l + v$, finding the unknown parameters x that minimize $v^T P v$, where P is the weight matrix reflecting the precisions and correlations of the observations [PEER-REVIEWED].

When CODATA encounters highly discrepant input data—such as the historical discrepancies in the Newtonian gravitational constant G or the muon magnetic-moment anomaly—the task group does not blindly apply the algorithm. The contradiction rule dictates that the *a priori* assigned uncertainties of the conflicting measurements are multiplied by an expansion factor such that each of their residuals from the adjustment falls at or below a threshold of 2 [PEER-REVIEWED]. This explicitly forces the model to treat the tension as a symptom of underestimated systematic error rather than statistical fluctuation.

Furthermore, the 2018 redefinition of the International System of Units (SI) demonstrates a case of synthesis fundamentally altering the underlying axioms of a field. By fixing exact values for the Planck constant h, the elementary charge e, the Boltzmann constant k, and the Avogadro constant N_A, CODATA effectively shifted these parameters from being variables adjusted by the LSA to immutable constants [PEER-REVIEWED]. This eliminated the need to consider vast amounts of historical data related to their determination, fundamentally restructuring the correlation matrix and drastically reducing uncertainties across dependent constants [PEER-REVIEWED]. The prototype of the kilogram and the triple point of water (TTPW) were retired as foundational inputs because they were no longer deemed fundamental [PEER-REVIEWED].

Mapping and Disanalogy: CODATA’s expansion of discrepancies maps to the kernel's Invariant I7 (transformation store with mandatory assumption manifests). By expanding the uncertainty, CODATA forces a re-evaluation of the assumption that an experiment's stated error budget is complete. The 2018 SI redefinition maps perfectly to Invariant I2 (append-only supersession). The entire claim graph was superseded by a new axiomatic foundation, rendering older derivation paths obsolete but preserving them in the audit footprint. The LSA’s reliance on minimizing residuals over continuous variables does not directly map to a typed logic graph, but the meta-procedure of forcing assumption audits upon encountering persistent residuals is isomorphic to querying the transformation store for shared failure points.

2.3 FLAG: Explicit Grading Rubrics for Lattice QCD

The Flavour Lattice Averaging Group (FLAG) reviews lattice Quantum Chromodynamics (QCD) results. Unlike the PDG or CODATA, the inputs for FLAG are completely theoretical computations rather than sensor data [PEER-REVIEWED]. Consequently, FLAG's synthesis machinery acts as a strict quality-control gatekeeper for theoretical frameworks.

FLAG utilizes an explicit, color-coded grading rubric to assess calculations. The criteria evaluate the quality of the continuum extrapolation, chiral extrapolation, and finite volume effects

[PEER-REVIEWED]. The grading determines whether a claim is granted entry into the global average:

- **Green Star ($\bigstar$):** The relevant systematic error has been estimated in a satisfactory manner and is under strict theoretical control. For example, for chiral extrapolation, a green rating requires the minimum pion mass $M_{\pi, min}$ simulated on the lattice to be less than 250 MeV [PEER-REVIEWED].
- **Empty Green Circle ($\circ$):** A reasonable attempt at estimating the systematic error has been made, but further refinement is necessary. For chiral extrapolation, this corresponds to $250 \text{ MeV} \le M_{\pi, min} \le 400 \text{ MeV}$ [PEER-REVIEWED].
- **Red Square ($\blacksquare$):** The calculation failed the quality criteria (e.g., $M_{\pi, min} > 400 \text{ MeV}$). The result is strictly dropped from FLAG averages [PEER-REVIEWED].

FLAG also organizes averages based on the number of dynamical quark flavors ($N_f = 2, 2+1, \text{ or } 2+1+1$) used in the simulations [PEER-REVIEWED]. Results referring to different N_f are never averaged together, as they correspond to different field-theoretic approximations of QCD [PEER-REVIEWED].

Mapping and Disanalogy: FLAG is the perfect precedent for Invariant I1 (typed graded claims). FLAG proves that theory-side synthesis requires explicit metadata tagging of the methodologies used to derive a claim. The rubric rendering acts as a strict entry condition for the synthesis. Because both FLAG and the physics-kernel synthesize theoretical derivations rather than sensor data, there is almost zero disanalogy here. The kernel can seamlessly adopt a grading lattice that gates entry into the transformation store based on the rigor of the LLM's derivation steps.

2.4 Nuclear Data Evaluation: ENDF and EXFOR

The Evaluated Nuclear Data File (ENDF) system, alongside the EXchange FORmat (EXFOR) database, treats the evaluation of cross-sections and resonance parameters as a standing profession [PEER-REVIEWED]. ENDF's critical synthesis mechanism is the inclusion of sophisticated covariance files to map uncertainties.

The evaluation process generates specific files: File 32 for resonance parameter covariances and File 33 for energy-dependent cross-section covariances [PEER-REVIEWED]. These files provide typed uncertainty provenance. They capture "long-range" covariance components that affect uncertainties over many neighboring resonances, as well as "short-range" components specific to individual interactions [PEER-REVIEWED]. Processing codes like PUFF-III and PUFF-IV extract these covariances, evaluating the eigenvalues of the correlation matrices and testing for positive definiteness before the data can be utilized in neutronics simulations [PEER-REVIEWED].

Mapping and Disanalogy: ENDF mechanizes Invariant I5 (computable provenance and audit footprints). It treats the uncertainty of a claim not as an isolated scalar value, but as a matrix that points backward to the specific experimental or theoretical assumptions that generated it. The kernel can adopt this by enforcing that all claims carry a covariance-like dependency matrix linking them to their assumption manifests (I7), enabling rapid identification of cascading errors when a foundational assumption is updated.

2.5 Astronomical Catalogs: Cross-Match Provenance

In multi-wavelength astronomy, cross-matching pipelines (such as the C3 Command-line Catalogue Cross-match tool or VizieR) align massive sky catalogs [PEER-REVIEWED]. The

procedure involves projecting angular distances and evaluating overlaps against precise threshold radii.
To optimize the O(N^2) computational complexity of cross-matching billions of records, algorithms convert the angular haversine distance into a 3D-tree Euclidean distance, partitioning the sky into cells and evaluating targets within Moore's neighborhood (the nine surrounding cells) [PEER-REVIEWED]. The synthesis generates "cross-neighbour" tables that explicitly record the provenance of a match, maintaining the identifiers of the masterObjID (the central source) and the slaveObjID (the identified counterpart from the external catalog) [PEER-REVIEWED].
Mapping and Disanalogy: This directly maps to Invariant I4 (first-class gaps with closing conditions). Cross-matching algorithms computationally execute a closing condition: if the coordinates and parameters of entity A and entity B fall within threshold \theta, the gap is closed, and they are synthesized as a single physical entity while preserving the origin pointers.

Institution	Synthesis Focus	Grading / Inclusion Rule	Contradiction Rule	Retirement Rule	Kernel Invariant
PDG	Particle Properties	Inverse-variance weights	Expand error by S = \sqrt{\chi^2 / N_{eff}}	Moved to obsolete archives	I6 (Held contradictions)
CODATA	Fundamental Constants	LSA over correlation matrix	Expand *a priori* errors until residuals \le 2	Superseded by definitional axioms (e.g., 2018 SI)	I7 (Assumption manifests), I2
FLAG	Lattice QCD Theory	Explicit color-coded rubrics	Excluded if rubric fails (Red square)	Superseded by newer N_f simulations	I1 (Typed graded claims)
ENDF	Nuclear Cross-sections	Positive definite covariance	Processed via File 32/33 formats	Replaced by new ENDF versions	I5 (Computable provenance)
Astro	Sky Catalogs	Threshold radius overlap	Flagged as separate objects	Superseded by higher-res surveys	I4 (Gaps and closing conditions)

3. Theory-Side Machinery: Synthesizing Without Averages

When synthesizing theoretical physics frameworks, there is often no numerical data to average. Physics instead synthesizes theory programs through topological bounds, logical constraints, and exclusion grids.

3.1 The Swampland Program as an Informal Claim Graph

The Swampland program seeks to identify low-energy effective field theories (EFTs) that appear semiclassically consistent but cannot be completed into quantum gravity in the ultraviolet regime [PEER-REVIEWED]. The program is synthesized via a "conjecture web."
Reviews by Palti and van Beest et al. explicitly structure the field as a logical network [PEER-REVIEWED]. The conjectures (e.g., the Swampland Distance Conjecture, the Weak

Gravity Conjecture, the No Global Symmetries Conjecture) act as nodes. Rigorous string theory derivations, holography arguments, and black hole entropy limits act as implication edges, connecting the nodes (e.g., showing how the Distance Conjecture implies the Weak Gravity Conjecture in specific limits) [PEER-REVIEWED]. Counterexamples act as refinement triggers; for instance, the discovery of specific Calabi-Yau compactifications acts as a counterexample that forces a refinement or supersession of a conjecture, bounding the parameter space of allowable EFTs [PEER-REVIEWED].
Assessment: The claim that the Swampland web is already an informal typed claim graph is **CONFIRMED**. The literature explicitly encodes directional implications and records counterexamples as typed boundaries on parameter space. The kernel can directly parse such a structure as a series of claims (I1) interconnected by dependency edges (I5), utilizing counterexamples as first-class kills (I3).

3.2 No-Go Theorems as First-Class Kills (Invariant I3)

Historical no-go theorems serve as the ultimate exclusion reservoirs in theory synthesis. The field synthesizes these by treating them as barriers with explicitly stated evasion conditions.

- **Coleman-Mandula Theorem (1967):** Proved that spacetime and internal symmetries can only combine trivially, meaning the symmetry group of the S-matrix must be a direct product of the Poincaré group and an internal symmetry group [PEER-REVIEWED]. The proof explicitly relied on named assumptions: the existence of a Poincaré subgroup, positive definite mass, a finite number of particles below an energy threshold, and an analytic S-matrix [PEER-REVIEWED].
- **Weinberg-Witten Theorem (1980):** Forbade the existence of massless particles with spin $j>1/2$ carrying a Lorentz-covariant current, and spin $j>1$ carrying a Lorentz-covariant stress-energy tensor, implying the graviton cannot be a composite particle in a standard QFT [PEER-REVIEWED]. The proof evaluated matrix elements of the charge and current for one-particle asymptotic states in the limit of zero momentum [PEER-REVIEWED].

To evade Coleman-Mandula, the community must violate the assumption that the symmetry algebra is a standard Lie algebra; introducing Lie superalgebras (supersymmetry) safely evades the kill [PEER-REVIEWED]. To evade Weinberg-Witten for gravity, one invokes gauge symmetries or emergent gravity models where fundamental Poincaré covariant diffeomorphism symmetry is absent [PEER-REVIEWED].
Mapping and Disanalogy: This practice maps verbatim to Invariant I3 (first-class kills with reinstatement conditions). The no-go theorem is the "kill." The named physical assumptions are the "reinstatement conditions." If a new branch of the transformation store explicitly violates Assumption 3 (e.g., utilizing a non-Lie algebra), the kill is computationally deactivated for that branch. The disanalogy is effectively zero; this is a pure topological mapping of logical bounds that the kernel can adopt instantly to manage its exclusion reservoir.

3.3 Dynamic Synthesis: Living Reviews and Community Processes

Living Reviews in Relativity pioneered the dynamic synthesis model, wherein review articles are continuously updated rather than static [PEER-REVIEWED]. Updates are typically triggered by major observational breakthroughs (e.g., the detection of high-frequency gravitational waves or new pulsar timing arrays) or significant theoretical shifts [PEER-REVIEWED]. Meanwhile, processes like Snowmass and the European Strategy synthesize community consensus to allocate funding and prioritize specific collider or detector technologies.

Mapping and Disanalogy: Living Reviews map securely to Invariant I8 (forkable recomputable standing), ensuring that a framework's state is perpetually current as new data enters. However, Snowmass processes present a fatal disanalogy: they are sociological and political mechanisms for attention allocation across thousands of physicists. The kernel, operating under one director, requires deterministic algorithmic allocation (e.g., prioritizing computational resources toward identified structural cruxes) rather than political consensus.

3.4 Executable Synthesis: The GAMBIT Global Fit Architecture

The Global And Modular BSM Inference Tool (GAMBIT) synthesizes beyond-Standard-Model parameter spaces [PEER-REVIEWED]. It is an executable synthesis framework that takes a Lagrangian, computes physical observables via various backend codes (e.g., Pythia, MicrOMEGAs, SPheno), and applies sophisticated statistical sampling (MCMC, nested sampling, differential evolution) to produce a composite likelihood function over the entire multi-dimensional parameter space of the model [PEER-REVIEWED].
Mapping and Disanalogy: GAMBIT relies on a combination of Invariant I9 (cheap structural queries) and I7 (transformation store). It demonstrates that to synthesize a theory program, the infrastructure must automatically resolve dependencies—determining which observable requires which backend code and executing them in the correct optimized order [PEER-REVIEWED]. The physics-kernel can emulate this topology by treating every theoretical claim as a computable node that dynamically updates a global state when new evidence is appended (I2). However, GAMBIT relies on executable numerical simulations, whereas the kernel currently synthesizes logical and semantic claims.

4. Tension Holding: Productive Contradictions

Physics rarely discards models instantly upon encountering a contradiction. The discipline holds them in tension, utilizing the resulting friction to isolate fundamental errors or discover new physics.

4.1 Muon g-2: Contradiction Between Synthesis Methods

The anomalous magnetic moment of the muon, $a_\mu = (g-2)/2$, currently exhibits a profound, multi-layered contradiction [PEER-REVIEWED]. The Muon g-2 Theory Initiative generated a Standard Model prediction utilizing a dispersive (data-driven) approach to calculate the Hadronic Vacuum Polarization (HVP) contribution [PEER-REVIEWED]. This method relies on the R-ratio extracted from extensive $e^+e^- \to \text{hadrons}$ cross-section data [PEER-REVIEWED]. This dispersive prediction sits at a $>5\sigma$ tension with Fermilab experimental measurements [PEER-REVIEWED].
However, a breakthrough lattice QCD calculation by the BMW collaboration generated an HVP prediction much closer to the experimental value, reducing the experiment-theory tension to $\sim 0.9\sigma$ [PEER-REVIEWED]. This created a new $\sim 4.0\sigma$ tension *between the two theoretical synthesis methods themselves* [PEER-REVIEWED]. Furthermore, new experimental data from CMD-3 created tensions within the e^+e^- datasets, undermining the dispersive consensus [PEER-REVIEWED].
The community manages this paralyzing contradiction by treating it as a load-bearing structure. They identified the "intermediate Euclidean time window" as the computational crux of the lattice

calculations [PEER-REVIEWED]. By computing only this specific, restricted window, independent lattice groups can rapidly verify or refute the BMW result without running the full, computationally exhaustive HVP calculation [PEER-REVIEWED].
Mapping and Disanalogy: This maps perfectly to Invariant I6 (held contradictions with computed cruxes). The kernel can mimic this behavior by automatically identifying the narrowest computational node (the equivalent of the "Euclidean window") where two divergent logical branches intersect, outputting it as the cheapest targeted query designed to resolve the overarching tension.

4.2 The Proton Radius Puzzle: Resolution Retrospectives

In 2010, the CREMA collaboration measured the proton charge radius via the Lamb shift in muonic hydrogen spectroscopy, yielding a highly precise result ($R_p = 0.8409 \pm 0.0004$ fm) that was completely incompatible with the 2010 CODATA value derived from electron scattering and standard hydrogen spectroscopy [PEER-REVIEWED]. The tension paralyzed the field for a decade.
The resolution emerged via the PRad experiment at JLab, which utilized a windowless hydrogen gas target and a magnetic-spectrometer-free method to measure at very small forward-scattering angles [PEER-REVIEWED]. This eliminated the unstated systematic errors present in previous electron scattering data (specifically, background scattering from the target windows).
What made this productive was the isolation of the contamination cone: the realization that the discrepancy was not due to new physics (e.g., lepton universality violation), but rather that a massive swath of standard electron scattering data shared a specific methodological vulnerability.
Mapping and Disanalogy: This maps directly to Invariant I5 (computable provenance and audit footprints). By tracing the provenance of the standard hydrogen scattering data, the kernel could explicitly define the contamination cone of the unstated systematic error, highlighting the exact experimental assumptions that required falsification.

4.3 Hubble Tension and Branch-Conditional Dependencies

The H_0 tension—the persistent discrepancy between early-universe measurements (CMB) and local distance ladder measurements (Cepheids/Supernovae)—is continuously quantified using Bayesian evidence and marginalized posteriors [FOLKLORE]. The literature taxonomy systematically organizes proposed resolutions into early-universe modifications (e.g., early dark energy) and late-universe modifications (e.g., void backreaction or separate universe conjectures) [PEER-REVIEWED]. The kernel would map this as a strict branch-conditional dependency, where the truth value of a cosmological claim is explicitly typed by its adherence to either the early or late universe resolution branch, preventing the silent conflation of incompatible physics.

5. Assumptions Discipline: The Theory Error Problem

In theoretical physics, underestimating "theory error" leads to phantom discoveries and false tensions. The rigorous tracking of assumptions is as critical as the tracking of experimental data.

5.1 Systematic Error Budgets and Blinding

Precision experimental physics relies heavily on systematic error budgets, where each systematic uncertainty is treated as a typed, provenanced contribution to the final error. To prevent confirmation bias—specifically the psychological tendency to tune analysis cuts until a result matches the Standard Model expectation—collaborations employ strict blinding and unblinding protocols [FOLKLORE]. This procedural discipline is functionally isomorphic to the kernel's generate-and-verify split. The generator (the LLM proposing a mechanism) operates blind to the downstream rigorous verification criteria, while the verifier assesses the claim against strict I1 rubrics without access to the generator's motivational context.

5.2 Documented Failures: Shared Assumptions in V_{ud} and PDFs

The most profound recent example of theory error resulting from unstated shared assumptions lies in the extraction of the Cabibbo-Kobayashi-Maskawa (CKM) matrix element V_{ud}. Superallowed $0^+ \to 0^+$ nuclear beta decays provide the most precise determination of V_{ud} [PEER-REVIEWED]. The extraction relies on correcting the experimental ft-values via a master formula that includes δ_C (isospin-symmetry breaking) and δ_{NS} (nuclear structure) corrections [PEER-REVIEWED].

For decades, these corrections were heavily reliant on the phenomenological shell-model calculations developed by Towner and Hardy [PEER-REVIEWED]. Because the entire community utilized these specific shared inputs, the theoretical uncertainty was perceived to be exceedingly small. However, recent re-evaluations using *ab initio* No Core Shell Models and dispersion relations revealed that the Towner-Hardy models missed critical effects, such as the contribution from the quasi-elastic nucleon absorption peak [PEER-REVIEWED]. Correcting this shared assumption induced a significant downward shift in V_{ud}, violating top-row CKM unitarity ($\Delta_{CKM} < 0$) and creating the "Cabibbo angle anomaly" [PEER-REVIEWED].

A similar vulnerability exists in the generation of Parton Distribution Functions (PDFs), which dictate how a proton's momentum is shared among its constituents. Global fits (e.g., CTEQ, MSTW, NNPDF) rely on diagonalizing Hessian matrices to determine error PDFs [PEER-REVIEWED]. Discrepancies often arise because different fitting groups treat correlated systematic errors in the underlying experimental data (like Tevatron jet cross-sections) differently, inadvertently baking shared historical biases into high-precision LHC predictions [PEER-REVIEWED].

Mapping and Disanalogy: These historical failures strongly validate Invariants I7 (mandatory assumption manifests) and I5 (computable shared ancestry). If the physics community possessed a typed claim graph, querying it would have instantly revealed that dozens of seemingly "independent" V_{ud} extractions and PDF error bounds ultimately terminated at single nodes: the Towner-Hardy phenomenological model or a specific treatment of a Tevatron systematic error. The kernel mechanizes this defense natively, preventing the illusion of independent verification when claims share a hidden foundational ancestor.

6. Synthesis Products the Kernel Could Emit

Based on the dissection of physics-native institutions, the following kernel-native synthesis products are evaluated.

Precedent	Preconditions	Earliest Failure	Cost Class	Verdict
No-Go Registry (Coleman-Mandula, Weinberg-Witten, Swampland bounds)	I3 (First-class kills) and I7 (Assumption manifests) fully populated and explicitly linked.	The stateless LLM hallucinates an evasion condition that violates basic topological logic, creating a false-positive reinstatement of a dead theory.	Low (Pure structural read over existing I3 nodes).	**CONFIRMED**. Pure topological mapping; no community disanalogy. Directly feeds Stage 3c reservoir design by rendering kills exactly as the field renders no-go theorems.
FLAG-style Rubric Rendering (Lattice QCD explicit grading)	I1 (Graded claims) with strict metadata attached to every theoretical derivation pathway.	The grading ontology becomes too rigid, failing to accommodate novel derivation methods (e.g., non-perturbative dualities vs standard perturbative limits), resulting in valid claims being dropped.	Low (Metadata filtering, tagging, and visual rendering).	**CONFIRMED**. Essential for maintaining graph hygiene and preventing heuristic LLM leaps from corrupting the core transformation store.
Tension Dashboard (Muon g-2 intermediate window, Hubble tension taxonomy)	I6 (Held contradictions) linked with I5 (Provenance). Requires logic capable of intersecting divergent graphs.	The LLM fails to compute a mathematically or physically valid "crux" (e.g., hallucinating an intermediate state that is unobservable in principle).	Medium (Requires graph intersection algorithms and careful prompt tuning).	**PLAUSIBLE**. High value for resolving state conflicts, but requires intense oversight to ensure the generated crux is a physically executable query.
PDG-style Scale Factor Generation (S = \sqrt{\chi^2 / N_{eff}})	Statistical variance data on contradictory claims.	Applying statistical variance scaling to LLM outputs. LLM systemic biases are highly correlated; applying a Birge ratio will mathematically launder correlated hallucinations into	High (Requires implementation of pseudo-statistical algorithms over text).	**DEAD**. The central disanalogy is fatal. The kernel must use topological cruxes to resolve contradictions, not statistical smoothing.

Precedent	Preconditions	Earliest Failure	Cost Class	Verdict
		false certainty.		
Global-Fit Consistency Report (GAMBIT architecture)	I9 (Structural queries) and I7 (Transformation store) mapped to a continuous parameter space.	The framework lacks the executable backends (Pythia, MicrOMEGAs) required to actually compute the physical likelihoods, reducing the report to qualitative, non-rigorous guessing.	Extreme (Requires building executable numerical interfaces into the graph).	**NEEDS REFINEMENT**. The kernel currently manages a *logical* claim graph, not an executable numerical simulation. It cannot compute continuous likelihoods without external integration.

7. Free Experiments: Executable Reads Over the Graph

Given a corpus of roughly 200 claims, 30 mechanisms, a 2-branch crux, and a 60-entry exclusion reservoir, the following pure reads over the existing graph can be executed by a session now, requiring no new infrastructure.

Experiment 1: The No-Go Registry Generator

- **Query Mechanism:** MATCH (k:Kill) -[:HAS_ASSUMPTION]-> (a:Assumption) RETURN k.name, a.name, k.reinstatement_condition
- **Precedent Procedure:** Modeled on the historical tracking of the Weinberg-Witten and Coleman-Mandula theorems [PEER-REVIEWED], where the field catalogs barriers alongside the exact modifications required to bypass them.
- **Expected Output:** A rendered Markdown table listing all 60 items in the exclusion reservoir. For each row, the output dictates the exact theoretical assumption that triggered the kill and the hypothetical physics modification that would reinstate it.
    - *Example Output Row:* "Theorem: No-global-symmetries | Assumption: UV-complete quantum gravity | Evasion Condition: Prove the specific framework is merely an isolated EFT explicitly decoupled from gravitational limits."

Experiment 2: Contamination Cone Trace (Shared Assumption Audit)

- **Query Mechanism:** MATCH (c1:Claim), (c2:Claim) MATCH (c1)-[:DEPENDS_ON*]->(a:Assumption) MATCH (c2)-[:DEPENDS_ON*]->(a) WHERE c1 <> c2 RETURN a.name, collect(c1.name), collect(c2.name)
- **Precedent Procedure:** Modeled on the retrospective diagnosis of the V_{ud} Cabibbo angle anomaly, where seemingly independent precision measurements collapsed because they shared the Towner-Hardy nuclear structure corrections [PEER-REVIEWED].

- **Expected Output:** An "Assumption Vulnerability Report." The stateless LLM parses the graph to find nodes in the 200-claim corpus that appear to be independently verified but actually converge on a single mechanism or assumption deep in the dependency tree. It highlights the single points of failure in the displacement geometry framework, mirroring how shared systematic errors in parton distribution functions compromise global fits [PEER-REVIEWED].

Experiment 3: FLAG-Style Entry Matrix

- **Query Mechanism:** MATCH (c:Claim)-[:DERIVED_BY]->(m:Mechanism) RETURN c.name, m.methodology, c.grade ORDER BY c.grade DESC
- **Precedent Procedure:** Modeled on the Flavour Lattice Averaging Group's (FLAG) color-coded inclusion rubric, which gates entry into global averages based on rigorous extrapolation controls [PEER-REVIEWED].
- **Expected Output:** A tiered, color-coded list of the 200 claims.
  - Claims derived via strictly verifiable mathematical steps are tagged [GREEN / EXACT].
  - Claims derived via heuristic LLM leaps or phenomenological estimates are tagged [YELLOW / ESTIMATE].
  - Claims lacking a strict formal derivation pathway are tagged [RED / DROPPED] and are algorithmically excluded from downstream mechanism building until the gap (Invariant I4) is closed with a rigorous proof.

Works cited

1. A simple tool for weighted averaging of inconsistent data sets - arXiv, https://arxiv.org/html/2406.08293v3 2. PoS(FPCP2015)016, https://pos.sissa.it/248/016/pdf 3. A minimalistic and general weighted averaging method for inconsistent data - arXiv, https://arxiv.org/html/2406.08293v1 4. REVIEW OF PARTICLE PHYSICS* Particle Data Group - Fermilab | Technical Publications, https://lss.fnal.gov/archive/openaccess/openaccess_cpc_40_10_100001.pdf 5. Introduction, https://ccwww.kek.jp/pdg/2025/reviews/rpp2024-rev-rpp-intro.pdf 6. arXiv:2004.01219v2 [physics.data-an] 11 Jun 2020, https://arxiv.org/pdf/2004.01219 7. I(JP) = 0(0−) https://pdg.lbl.gov Page 1 Created: 6/1/2021 08:32, http://pdg.ge.infn.it/2021/listings/rpp2021-list-Ds-plus-minus.pdf 8. Precise measurement of the top-quark mass at the CMS experiment using the ideogram method - INIS-IAEA, https://inis.iaea.org/records/b08ay-6kd84/files/47001422.pdf?download=1 9. Review of Particle Properties, 1969 - SciSpace, https://scispace.com/pdf/review-of-particle-properties-1969-422vcor2nh.pdf 10. Review of Particle Properties Particle Data Group - SciSpace, https://scispace.com/pdf/review-of-particle-properties-particle-data-group-2nkc9qyioo.pdf 11. CODATA Recommended Values of the Fundamental Physical Constants: 2018, https://pml.nist.gov/cuu/pdf/JPCRD2018CODATA.pdf 12. CODATA recommended values of the fundamental physical constants: 2018 - PMC - NIH, https://pmc.ncbi.nlm.nih.gov/articles/PMC9890581/ 13. Least-squares adjustment - Grokipedia, https://grokipedia.com/page/Least-squares_adjustment 14. arXiv:1504.01633v1 [physics.data-an] 7 Apr 2015, https://arxiv.org/pdf/1504.01633 15. THE 2006 ADJUSTMENT OF THE FUNDAMENTAL CONSTANTS - National Institute of Standards and Technology,

https://tsapps.nist.gov/publication/get_pdf.cfm?pub_id=32919 16. The 1986 CODATA Recommended Values Of the Fundamental Physical Constants - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC5287000/ 17. arXiv:2409.03787v1 [hep-ph] 30 Aug 2024, https://arxiv.org/pdf/2409.03787 18. CODATA recommended values of the fundamental physical constants: 2022*, https://physics.nist.gov/cuu/pdf/RevModPhys.97.025002.pdf 19. 1 Introduction - FLAG, http://flag.itp.unibe.ch/2019/Media?action=AttachFile&do=get&target=FLAG_introduction.pdf 20. Flavour and Lattice, https://indico.in2p3.fr/event/6001/contributions/36438/attachments/29211/35999/moriond.pdf 21. FLAG Review 2019, http://flag.unibe.ch/2019/ 22. Assignment of uncertainties to scientific data - INIS-IAEA, https://inis.iaea.org/records/pr4q4-e6w02 23. ENDF-324 Vol.4 - National Nuclear Data Center, https://www.nndc.bnl.gov/endfdocs/ENDF-324-Vol4.pdf 24. Recent Advances with the AMPX Covariance Processing Capabilities in PUFF-IV - National Nuclear Data Center, https://www.nndc.bnl.gov/cw2008/presentations/wiarda.pdf 25. Processing of ENDF/B-VII Covariance Data for Use with Sensitivity/Uncertainty Analysis - ANS / Digital Nuclear Library, https://epubs.ans.org/download/?a=6221 26. The Data Covariance Files for ENDF/B-V OAK RIDGE NATIONAL LABORATORY - INIS-IAEA, https://inis.iaea.org/records/rqnv9-7c042/files/8345665.pdf?download=1 27. PUFF-III: A Code for Processing ENDF Uncertainty Data Into Multigroup Covariance Matrices-(NUREG/CR-6650)-003739315, https://www.nrc.gov/docs/ml0037/ML003739315.pdf 28. C³: A Command-line Catalogue Cross-matching tool for modern astrophysical survey data | Request PDF - ResearchGate, https://www.researchgate.net/publication/314283247_C_A_Command-line_Catalogue_Cross-matching_tool_for_modern_astrophysical_survey_data 29. Top 9 Best Astronomical Software: 2026 Comparison - Gitnux, https://gitnux.org/best/astronomical-software/ 30. Astroinformatics Tools - Cambridge University Press & Assessment, https://www.cambridge.org/core/services/aop-cambridge-core/content/view/90EDFB70DE348060BA84ECD3DD809EE6/S1743921316013120a.pdf/div-class-title-c-span-class-sup-3-span-a-command-line-catalogue-cross-matching-tool-for-modern-astrophysical-survey-data-div.pdf 31. A Low-complexity Cross-matching Algorithm for Large Astronomical Catalogues Using 3d-tree Approach - Research in Astronomy and Astrophysics (RAA), https://www.raa-journal.org/issues/all/2023/v23n10/202310/P020241029349196441523.pdf 32. SB9 catalogue: status, comparison with non-single stars from Gaia DR3, and evolution to SBX - Oxford Academic, https://academic.oup.com/mnras/article/547/2/stag351/8499613 33. An Introductory SQL Tutorial for Gaia-ESO Science Archive Users - GES, http://ges.roe.ac.uk/sqlcookbook_idr1.html 34. From Frame Covariance to the Swampland Distance Conjecture - arXiv, https://arxiv.org/pdf/2512.07929 35. arXiv:1903.06239v1 [hep-th] 14 Mar 2019, http://alpha.sinp.msu.ru/~panov/Lib/Papers/SWAMPLAND/1903.06239.pdf 36. Swampland (physics) - Grokipedia, https://grokipedia.com/page/Swampland_(physics) 37. Contents - arXiv, https://arxiv.org/html/2506.03253v2 38. Swampland Conjectures - TracingCurves, https://tracingcurves.wordpress.com/tag/swampland-conjectures/ 39. arXiv:2301.08262v3 [hep-th] 3 Apr 2023, https://arxiv.org/pdf/2301.08262 40. Coleman–Mandula theorem - Wikipedia, https://en.wikipedia.org/wiki/Coleman%E2%80%93Mandula_theorem 41. The Coleman-Mandula Theorem, https://www.mathi.uni-heidelberg.de/~walcher/teaching/sose16/geo_phys/ColemanMandula.pdf 42. Emerging Internal Symmetries from Effective Spacetimes - MPG.PuRe, https://pure.mpg.de/rest/items/item_2486842_2/component/file_2486841/content 43.

Coleman-Mandula Theorem Explained | PDF | Group (Mathematics) | Algebra - Scribd, https://www.scribd.com/document/877082704/Coleman-Mandula 44. Weinberg–Witten theorem - Wikipedia, https://en.wikipedia.org/wiki/Weinberg%E2%80%93Witten_theorem 45. Weinberg–Witten theorem - Grokipedia, https://grokipedia.com/page/Weinberg%E2%80%93Witten_theorem 46. A Study of Relativistic Bounds on Clock Synchronisation on Earth - Defence Science and Technology, https://www.dst.defence.gov.au/sites/default/files/publications/documents/DST-Group-RR-0454_0.pdf 47. Living Reviews in Relativity author index - Leo C. Stein, https://duetosymmetry.com/lrr-index/ 48. [1304.3473] Gravitational Wave Tests of General Relativity with Ground-Based Detectors and Pulsar Timing Arrays - arXiv, https://arxiv.org/abs/1304.3473 49. [2407.16540] Gravity experiments with radio pulsars - arXiv, https://arxiv.org/abs/2407.16540 50. GAMBIT - Indico, https://indico.cern.ch/event/875077/contributions/4489009/attachments/2297706/3908183/GAMBIT.pdf 51. Understanding dark matter with Gambit - Agenda INFN, https://agenda.infn.it/event/35353/contributions/237690/attachments/125931/185894/Dark%20matter%20with%20GAMBIT.pdf 52. The GAMBIT Universal Model Machine: from Lagrangians to likelihoods - DiVA Portal, https://www.diva-portal.org/smash/get/diva2:1626012/FULLTEXT01.pdf 53. Model Independent Tests of the Hadronic Vacuum Polarization Contribution to the Muon g-2 - arXiv, https://arxiv.org/html/2408.01123v2 54. Magnetic Moments of Leptons: What they reveal to us∗, https://people.physik.hu-berlin.de/~fjeger/Katowice25.pdf 55. (PDF) Muon g−2 and Δα connection - ResearchGate, https://www.researchgate.net/publication/343778762_Muon_g-2_and_Da_connection 56. arXiv:2408.01123v2 [hep-ph] 26 Mar 2025, https://arxiv.org/pdf/2408.01123 57. The anomalous magnetic moment of the muon: status of lattice QCD calculations - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC8023536/ 58. How I learnt to stop worrying and love QCD predictions - CERN Courier, https://cerncourier.com/how-i-learnt-to-stop-worrying-and-love-qcd-predictions/ 59. Peter Boyle High Energy Theory Physics @ BNL, https://www.orau.gov/support_files/2024scidac/breakout/hep/HEP-Peter-Boyle1.pdf 60. 2018 CAP Congress / Congrès de l'ACP 2018 - Indico Global, https://indico.global/event/427/timetable/?view=standard 61. Nature November | PDF | Science | Theory - Scribd, https://www.scribd.com/document/435800134/Nature-November 62. The importance of GR's principle of equivalence for kinematically determined universes and consequences for ΛCDM - Frontiers, https://www.frontiersin.org/journals/astronomy-and-space-sciences/articles/10.3389/fspas.2025.1698879/full 63. arXiv:2304.03800v2 [nucl-th] 24 Mar 2024, https://arxiv.org/pdf/2304.03800 64. Cabibbo angle anomalies and oblique corrections: The remarkable role of the vectorlike quark doublet - CERN, https://scoap3-prod-backend.s3.cern.ch/media/files/79625/10.1103/PhysRevD.108.035022.pdf 65. Toward ab-initio nuclear theory calculations of δC - ResearchGate, https://www.researchgate.net/publication/379488336_Toward_ab-initio_nuclear_theory_calculations_of_dC 66. Dispersive formalism for the nuclear structure correction - OSTI.GOV, https://www.osti.gov/servlets/purl/2419501 67. Superallowed nuclear beta decays and precision tests of the Standard Model - arXiv, https://arxiv.org/pdf/2311.00044 68. The form factor expansion in the precision β decay era - Frontiers, https://www.frontiersin.org/journals/physics/articles/10.3389/fphy.2026.1884614/pdf 69. High-precision determination of radiative corrections to superallowed nuclear beta decays - EPJ

Web of Conferences, https://www.epj-conferences.org/articles/epjconf/pdf/2025/27/epjconf_NuclearPhysics2025_01025.pdf 70. Quantum Monte Carlo calculation of δ_C in the superallowed beta decay of ^10C - arXiv, https://arxiv.org/pdf/2605.14006 71. arXiv:2112.10942v2 [hep-ph] 24 Feb 2022, https://arxiv.org/pdf/2112.10942 72. Taming nuclear size and shape effects in superallowed beta-decay - arXiv, https://arxiv.org/pdf/2605.13985 73. 1 Parton Distribution Functions, https://web.pa.msu.edu/people/huston/atlas/higgs_pdf/PDF_temp.pdf 74. Parton Distribution Functions | The Black Book of Quantum Chromodynamics: A Primer for the LHC Era | Oxford Academic, https://academic.oup.com/book/27902/chapter/203882272 75. Parton distribution functions - Indico, https://indico.bnl.gov/event/603/contributions/15145/attachments/13511/16557/02_d816.pdf